

Sheet 2

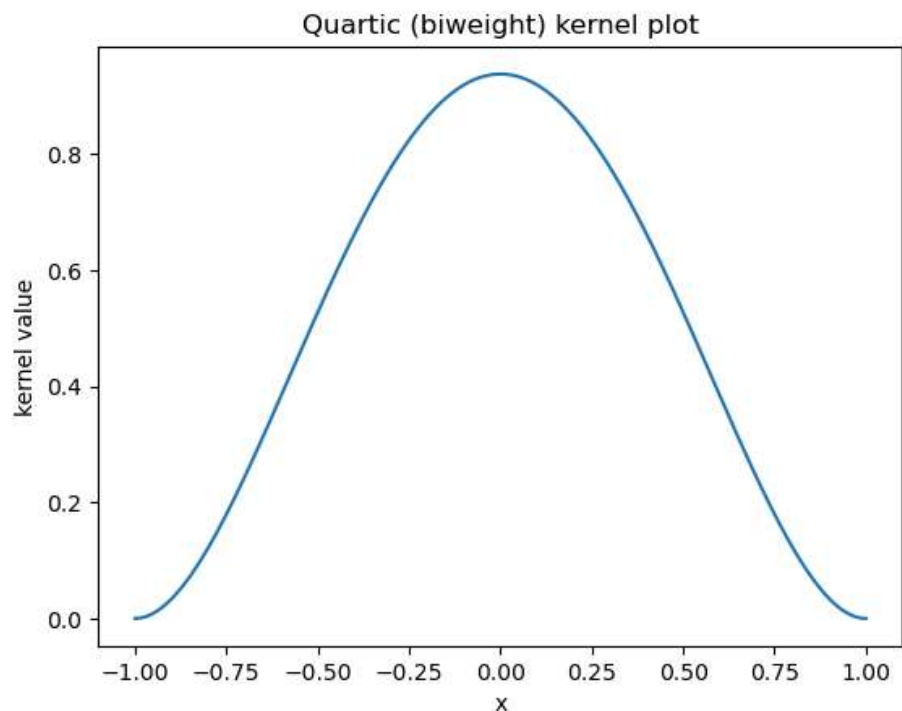
```
In [1]: import numpy as np
        from matplotlib import pyplot as plt
        from scipy.stats import gaussian_kde
```

1 Kernel Density Estimation

(a)

```
In [2]: def biweight(x, mu, w):
        """biweight kernel at mean mu, with bandwidth w evaluated
        if (abs(x - mu) <= w):
            return 15. * (1. - ((x - mu)/w)**2)**2 / (16. * w)
        else:
            return 0
        #TODO: implement the quartic (biweight) kernel
```

```
In [ ]: # TODO plot the kernel
        x = np.linspace(-1, 1, 1000)
        y = [biweight(t, mu=0, w=1) for t in x]
        fig, ax = plt.subplots()
        ax.set_xlabel("x")
        ax.set_ylabel("kernel value")
        ax.set_title("Quartic (biweight) kernel plot")
        ax.plot(x, y)
        plt.show()
```



(b)

```
In [ ]: # Load the data
data = np.load("data/samples.npy")
data50 = data[:50]
print(f'{data.shape=}, {data50.shape=}')
```

data.shape=(10000,), data50.shape=(50,)

```
In [15]: def kde(x, obs, w=1):
return sum(biweight(x, data_x, w) for data_x in obs) / len
# TODO: implement the KDE with the biweight kernel
```

```
In [16]: # TODO: compute and plot the kde on the first 50 data points
ws = [0.1, 0.5, 1, 2, 3, 5, 7]

fig, ax = plt.subplots(len(ws), 1, sharey='col')

for i, w in zip(range(len(ws)), ws):
    x = np.linspace(-10, 20, 1000)
    y = [kde(graph_x, data50, w) for graph_x in x]
    ax[i].plot(x, y)
    ax[i].set_xlabel('x')
    ax[i].set_ylabel('density')
    ax[i].set_title(f'Kernel density estimation with bandwidth {w}')

fig.set_figwidth(15)
fig.set_figheight(40)
plt.show()

print("Interpretation: The bandwidths 0.1 and 0.5 are clearly
      "because the influence of certain data points is too large
      "Plots with bandwidths 1, 2 and 3 are much smoother, big
      "Although with the increase of the bandwidth the slope decreases
      "Bandwidths 5 and 7 seem to be too large,\n"
      "in their respective graphs the underlying structure of the data is lost
      "The optimal bandwidth is w=2, it preserves the structure of the data
      "while being smooth enough to not have effects from single data points
      ")

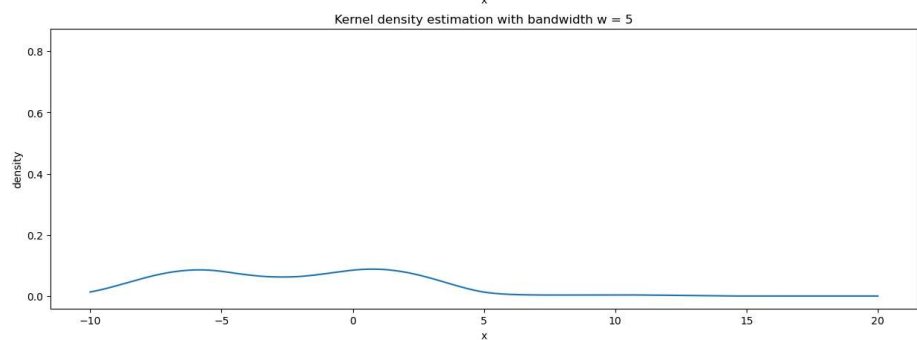
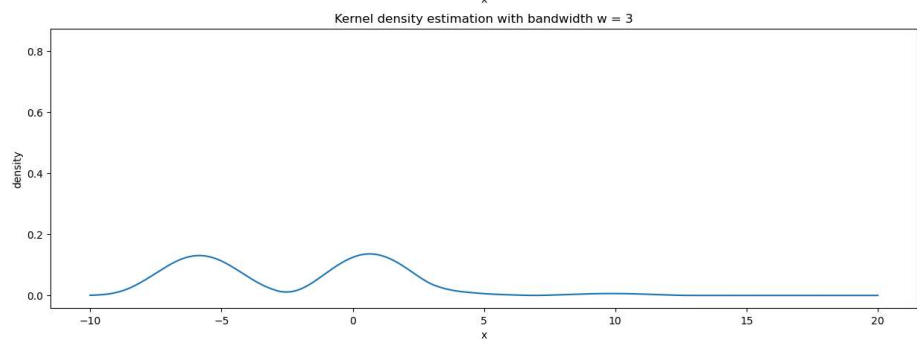
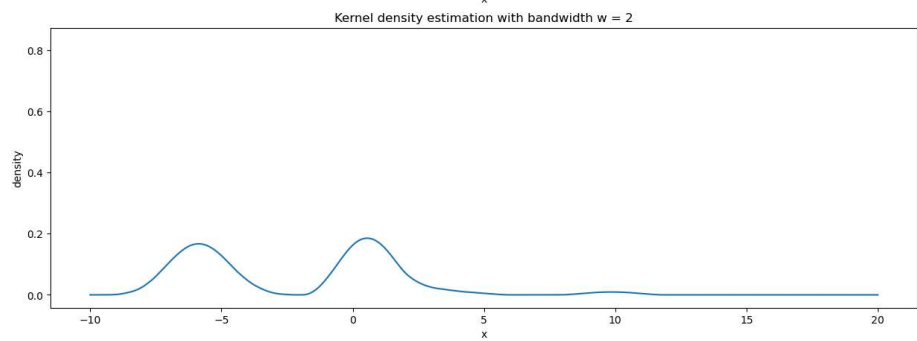
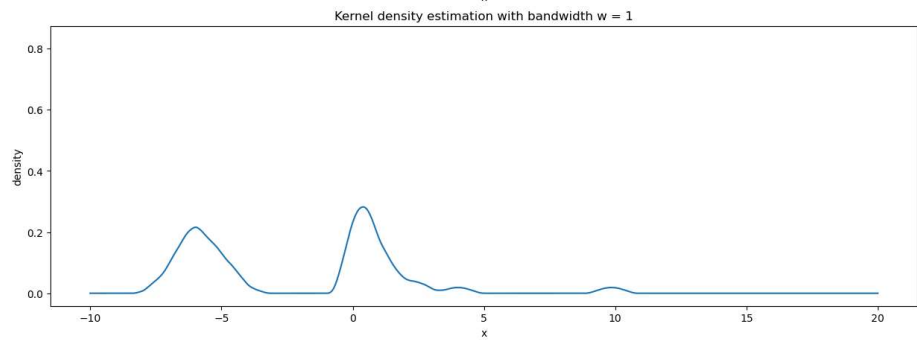
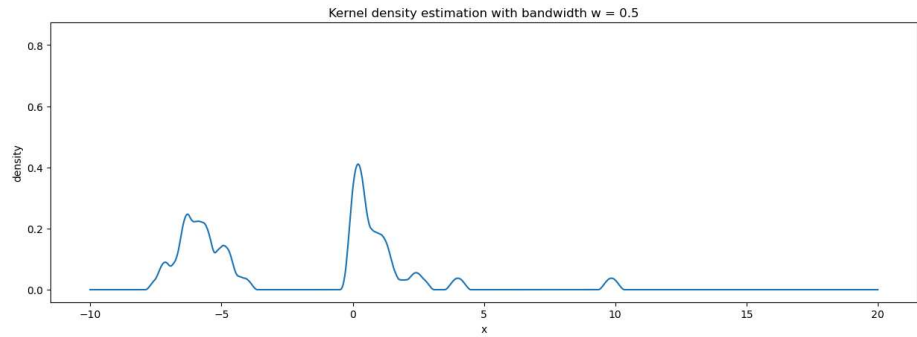
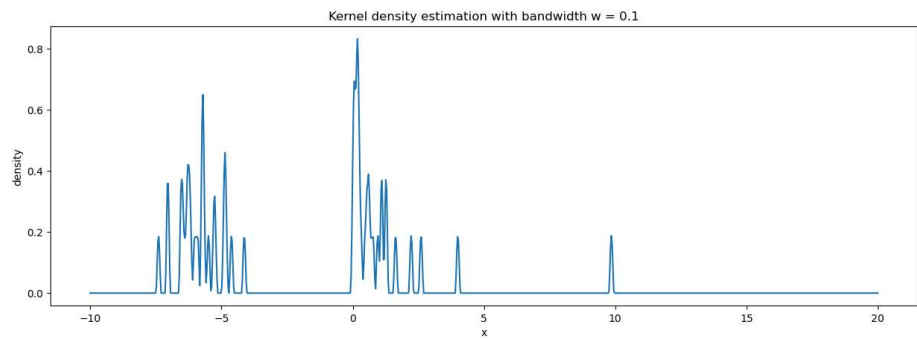
# TODO: explore what happens when you increase the number of points
data_len = [100, 1000, 10000]

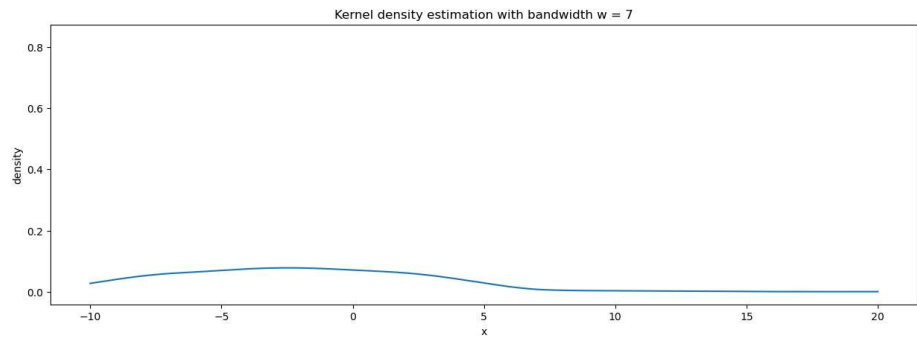
fig, ax = plt.subplots(len(ws), len(data_len), sharey='col')

for i, w in enumerate(ws):
    for j, N in enumerate(data_len):
        x = np.linspace(-10, 20, 100)
        data_curr = data[:N]
        y = [kde(graph_x, data_curr, w) for graph_x in x]
        ax[i, j].plot(x, y)
        ax[i, j].set_xlabel('x')
        ax[i, j].set_ylabel('density')
        ax[i, j].set_title(f'Kernel density estimation\n with {w} and {N} points')

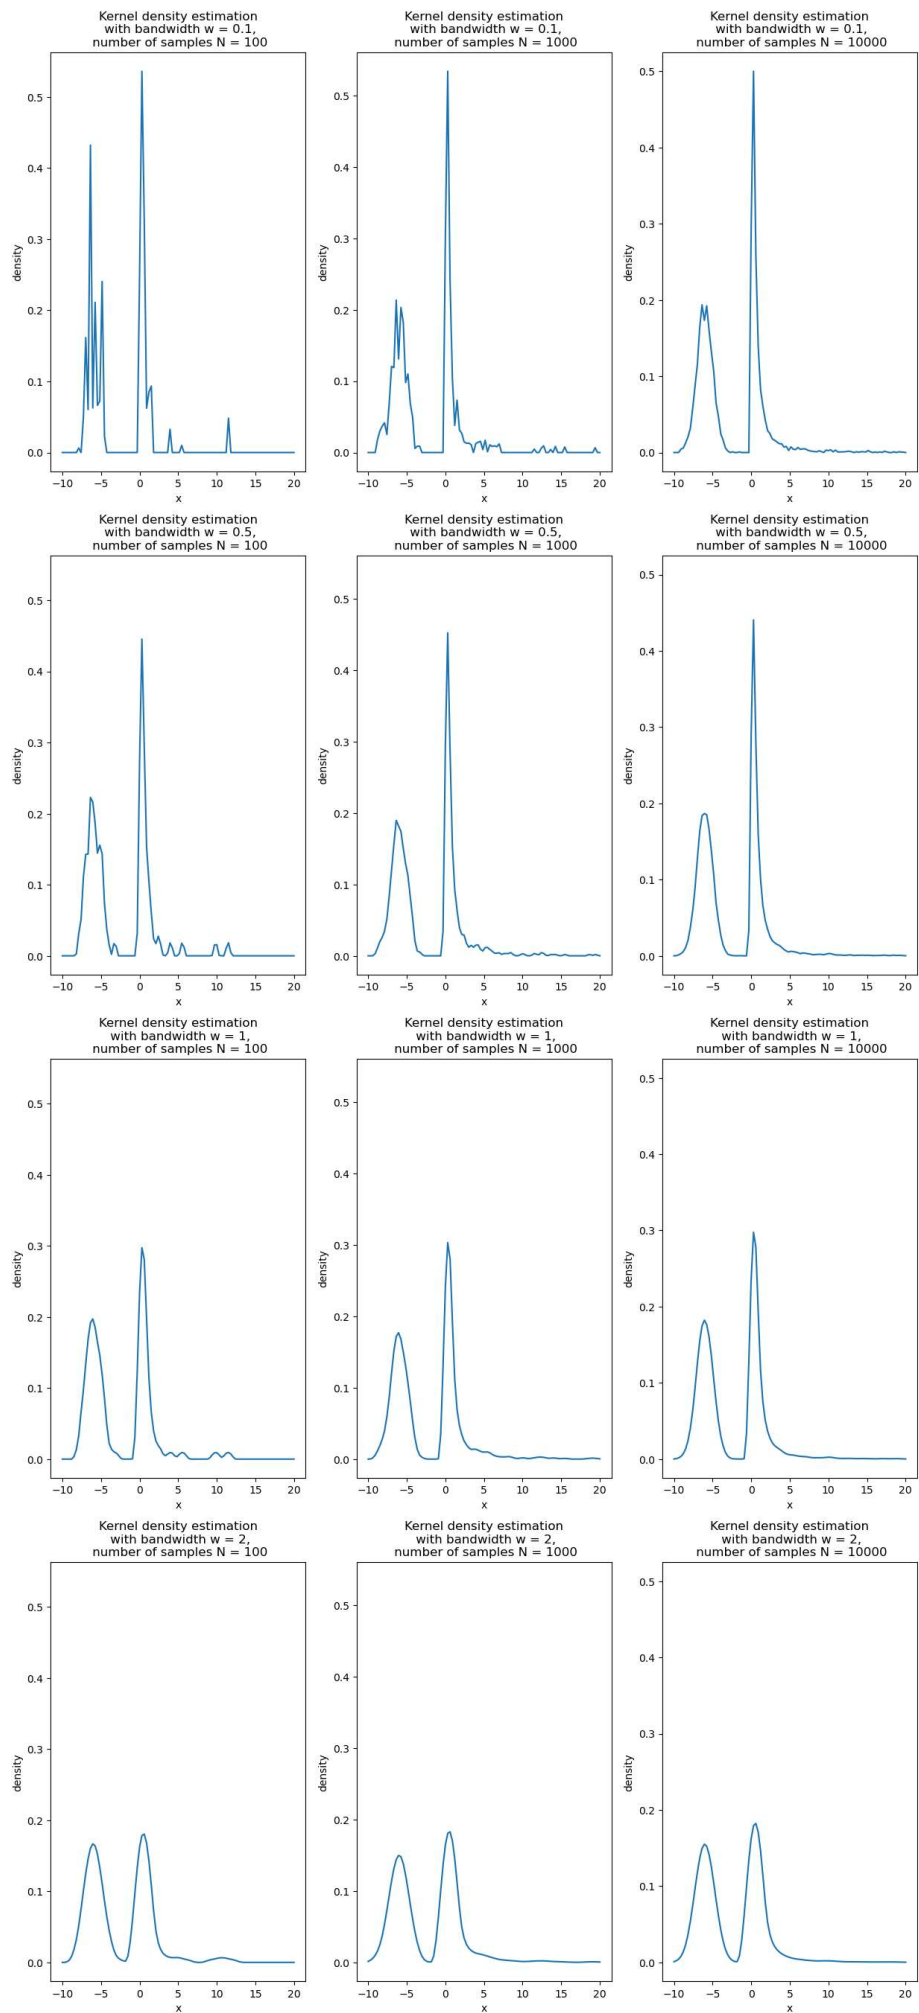
fig.set_figwidth(15)
fig.set_figheight(60)
plt.show()
```

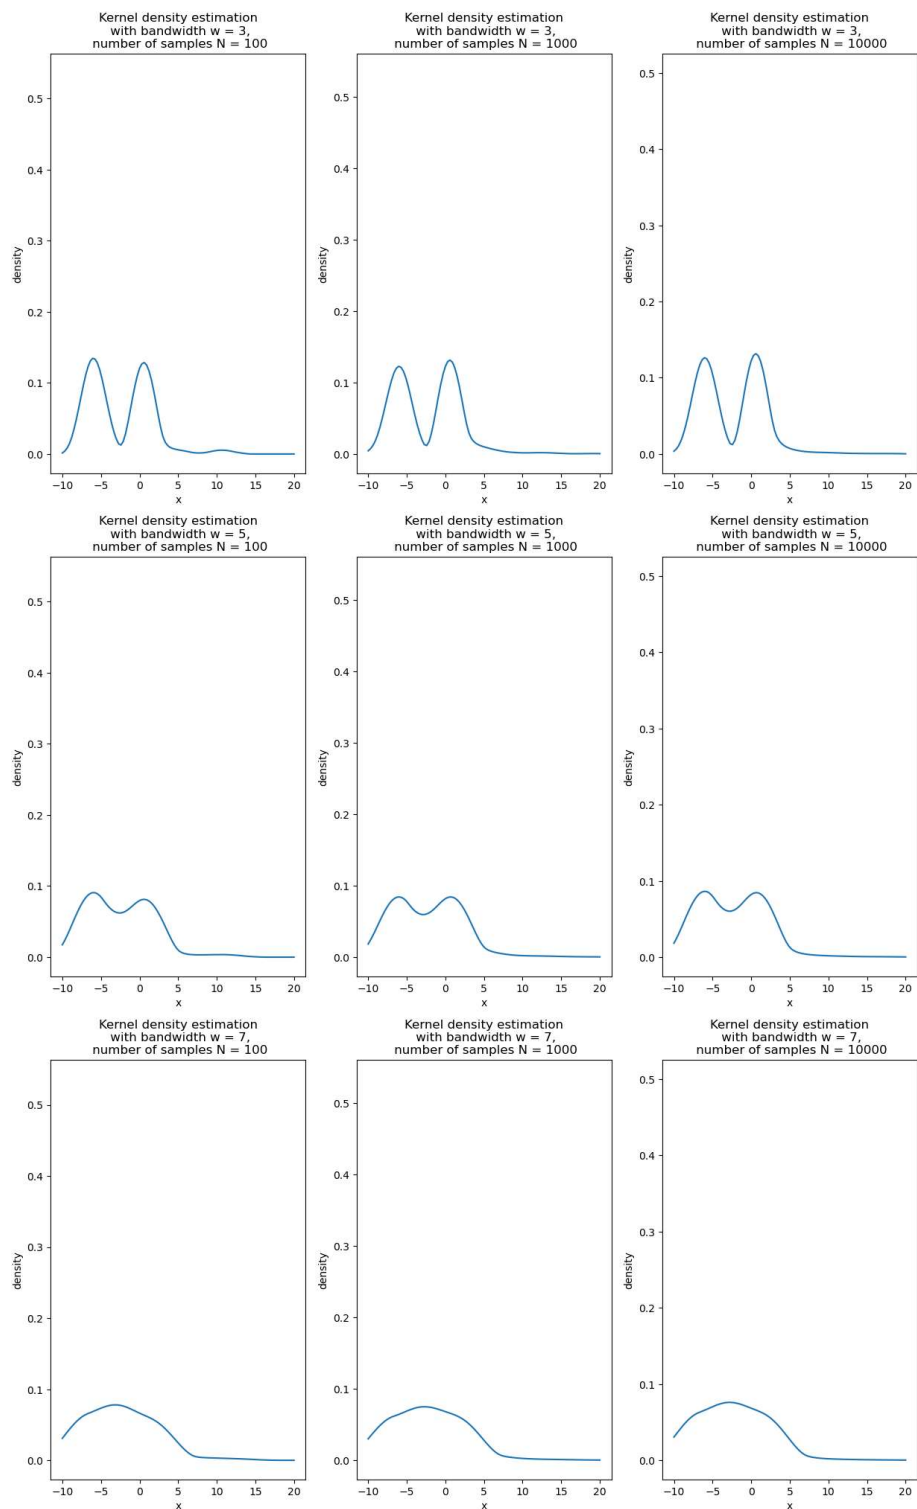
```
print("Interpretation: With the increase in the number of pairs  
      "For instance with the bandwidth  $w = 0.5$  graph with 100  
      "while the graph with 10000 samples with the same bandwidth  
      "Thus, a bigger number of data samples allows for a smaller  
      "On the bigger bandwidths the number of samples does not  
      ")
```





Interpretation: The bandwidths 0.1 and 0.5 are clearly too small since the graphs look too spiky, because the influence of certain data points is too large. Plots with bandwidths 1, 2 and 3 are much smoother, bigger bandwidths alleviate the effects of data artifacts. Although with the increase of the bandwidth the slope of the "bells" becomes less steep. Bandwidths 5 and 7 seem to be too large, in their respective graphs the underlying structure of the distribution (two bells) starts to disappear. The optimal bandwidth is $w=2$, it preserves the structure of the distribution and even the asymmetry of two bells, while being smooth enough to not have effects from singular data points.





Interpretation: With the increase in the number of points the graphs become smoother on smaller bandwidths.

For instance with the bandwidth $w = 0.5$ graph with 100 samples still has small spikes, while the graph with 10000 samples with the same bandwidth is almost completely smooth.

Thus, a bigger number of data samples allows for a smaller bandwidth, resulting in a more precise estimation.

On the bigger bandwidths the number of samples does not change the graph.

3 Mean-Shift

(b)

```
In [ ]: # For 3(a) see pdf
```

```
In [54]: # TODO: implement the update to the Local mean

def mean_shift_step(x, xt, r=1):
    """
    A single step of mean shift, moving every point in xt to the
    local mean of the points in x that are within distance r of it.

    Parameters
    -----
    x : np.ndarray
        Array of points underlying the KDE, shape (d, N1)
    xt : np.ndarray
        Current state of the mean shift algorithm, shape (d, N2)
    n_components : int, optional
        Number of requested components. By default returns all components.

    Returns
    -----
    np.ndarray
        the points after the mean-shift step
    """
    # NOTE: For the exercise you only need to implement this
    #       If you want some extra numpy-practice, implement it

    assert xt.shape[0] == x.shape[0], f'Shape mismatch: {x.shape[0]} vs {xt.shape[0]}'

    # TODO: start by computing a N by N matrix 'dist' of distances
    #       such that dists[i, j] is the distance between x[i] and xt[j]

    dist = np.zeros((x.shape[0], xt.shape[0]))

    for i, x_i in enumerate(x):
        for j, xt_j in enumerate(xt):
            dist[i][j] = np.linalg.norm(x_i - xt_j)

    local_means = np.zeros_like(xt)
    for j in range(len(xt)):
        sum_val = 0
        sum_num = 0
        for i in range(len(x)):
            if dist[i][j] < r:
                sum_val += x[i]
                sum_num += 1
        local_means[j] = sum_val / sum_num

    return local_means
```

```
In [ ]: # Load the data
data = np.load("data/samples.npy")
x = data[:200] # use e.g. the first 200 points
xt = x
```



```

trajectories = [xt]
max_steps = 100
for step in range(max_steps):

    # TODO: update xt with your mean shift step
    xt = mean_shift_step(x, xt)
    trajectories.append(xt)
    if np.allclose(trajectories[-1], trajectories[-2]): # break
        break
trajectories = np.stack(trajectories)
n_steps = len(trajectories) - 1

fig, ax = plt.subplots()

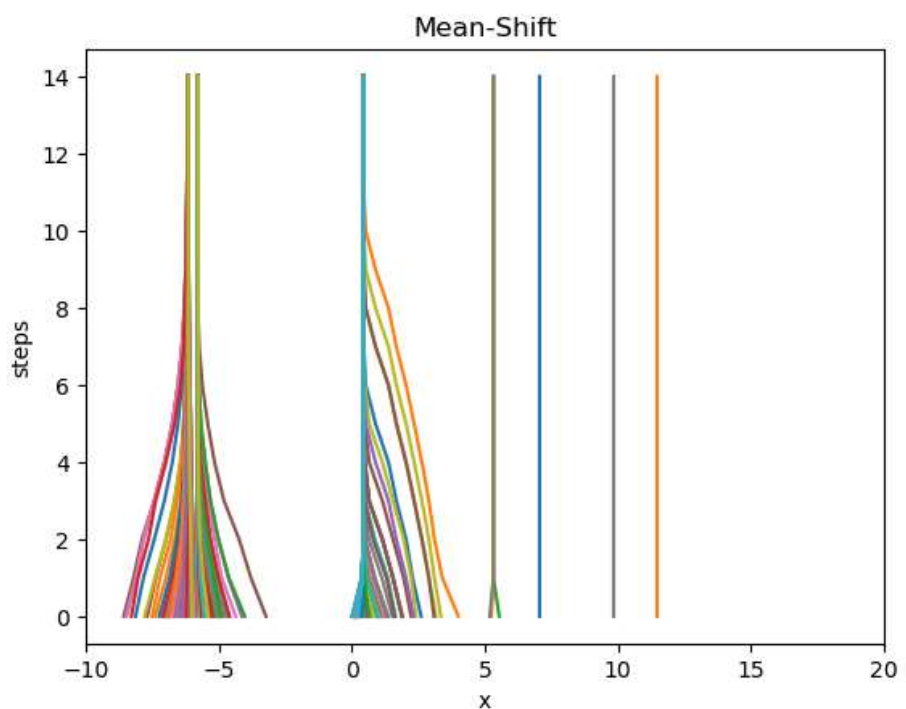
for i in range(trajectories.shape[1]):
    ax.plot(trajectories[:, i], np.array(range(trajectories.shape[0])))

ax.set_xlim(-10, 20)
ax.set_xlabel('x')
ax.set_ylabel('steps')
ax.set_title('Mean-Shift')

plt.show()

# TODO: plot the trajectories

```



```

In [ ]: # TODO: implement the update to the Local mean

def mean_shift_step_blurry(xt, r=1):
    """
    A single step of mean shift, moving every point in xt to the
    local mean.

    Parameters
    -----
    xt : np.ndarray
        Current state of the mean shift algorithm, shape (d, N)
    """

```

```

Returns
-----
np.ndarray
    the points after the mean-shift step
"""
# NOTE: For the exercise you only need to implement this
#       If you want some extra numpy-practice, implement it

# TODO: start by computing a N by N matrix 'dist' of distances
#       such that dists[i, j] is the distance between x[i] and x[j]

dist = np.zeros((xt.shape[0], xt.shape[0]))

for i, xt_i in enumerate(x):
    for j, xt_j in enumerate(xt):
        dist[i][j] = np.linalg.norm(xt_i - xt_j)

local_means = np.zeros_like(xt)
for j in range(len(xt)):
    sum_val = 0
    sum_num = 0
    for i in range(len(xt)):
        if dist[i][j] < r:
            sum_val += xt[i]
            sum_num += 1
    local_means[j] = sum_val / sum_num

return local_means

```

```

In [64]: # Load the data
data = np.load("data/samples.npy")
x = data[:200] # use e.g. the first 200 points
xt = x

trajectories = [xt]
max_steps = 100
for step in range(max_steps):
    # TODO: update xt with your mean shift step
    xt = mean_shift_step_blurry(xt)
    trajectories.append(xt)
    if np.allclose(trajectories[-1], trajectories[-2]): # break
        break
trajectories = np.stack(trajectories)
n_steps = len(trajectories) - 1

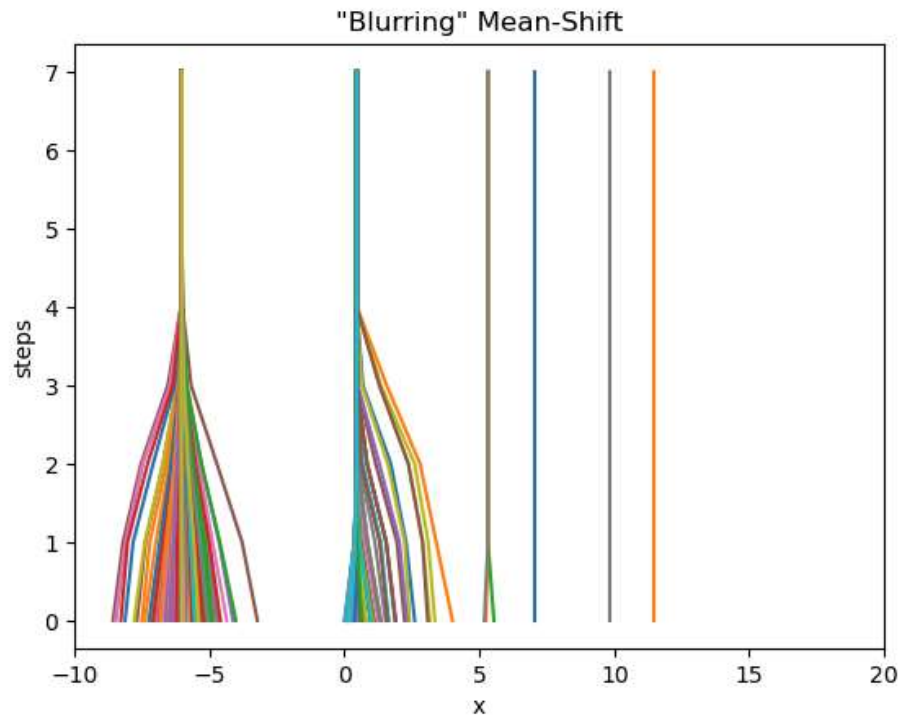
fig, ax = plt.subplots()

for i in range(trajectories.shape[1]):
    ax.plot(trajectories[:, i], np.array(range(trajectories.shape[0])))

ax.set_xlim(-10, 20)
ax.set_xlabel('x')
ax.set_ylabel('steps')
ax.set_title('\n"Blurring" Mean-Shift')
plt.show()

```

```
# TODO: plot the trajectories
```



Convergence of the "blurring" mean-shift is faster than the regular mean-shift. In the regular mean-shift the leftmost cluster was split into two, because it had two close local means, the blurring mean-shift, however, handled that and all the points in that cluster converged to one point.

Other than that the clusters are the same, we see two big clusters which coincide with the "bells" of the estimated density function and four small clusters which may be considered as noise/outliers.

3(a)

$$x_j^{t+1} = x_j^t + \alpha_j^t \frac{2}{n} \sum_{i: \|x_i - x_j^t\| < 1} (x_i - x_j^t) - \text{gradient ascent (1)}$$

local mean $\tilde{x}_j^{t+1} = \frac{\sum_{i: \|x_i - x_j^t\| < 1} x_i}{\sum_{i: \|x_i - x_j^t\| < 1} 1}$

$$\frac{\sum_{i: \|x_i - x_j^t\| < 1} x_i}{\sum_{i: \|x_i - x_j^t\| < 1} 1} = x_j^t + \alpha_j^t \frac{2}{n} \left(\sum_{i: \|x_i - x_j^t\| < 1} x_i - x_j^t \cdot \sum_{i: \|x_i - x_j^t\| < 1} 1 \right)$$

$$\alpha_j^t = \frac{\sum_{i: \|x_i - x_j^t\| < 1} x_i - x_j^t \cdot \sum_{i: \|x_i - x_j^t\| < 1} 1}{\sum_{i: \|x_i - x_j^t\| < 1} 1} \cdot \frac{1}{\sum_{i: \|x_i - x_j^t\| < 1} x_i - x_j^t \cdot \sum_{i: \|x_i - x_j^t\| < 1} 1} \cdot \frac{n}{2}$$

$$\alpha_j^t = \frac{n}{2 \cdot \sum_{i: \|x_i - x_j^t\| < 1} 1}$$

When substituting α_j^t with this fraction, the last term of nodes like this: $\frac{\sum_{i: \|x_i - x_j^t\| < 1} (x_i - x_j^t)}{\sum_{i: \|x_i - x_j^t\| < 1} 1}$ (1)

which is the mean vector of the neighbourhood - a sensible choice